

CLOUD COMPUTING

Session 3 :

High Availability and Scalability In The AWS Cloud

APROBAT

Conf. univ. dr. ing. IUSTIN PRIESCU - UTM
Dr. ing. Sebastian NICOLAESCU - Verizone Bussines-US

Design your application
assuming that:

“Everything fails, all the time”

Werner Vogel, PhD – Amazon CTO

What is High Availability?

High availability (HA) is about ensuring that your application's **downtime is minimized** as much as possible, without the need for human intervention.


Levels of Availability:

	Percent of Uptime	Max Downtime per Year	Equivalent Downtime per Day
1 Nine	90%	36.5 days	2.4 hrs
2 Nines	99%	3.65 days	14 min
3 Nines	99.9%	8.76 hrs	86 sec
4 Nines	99.99%	52.6 min	8.6 sec
5 Nines	99.999%	5.25 min	.86 sec

Best Practice: Avoid Single Points Of Failure

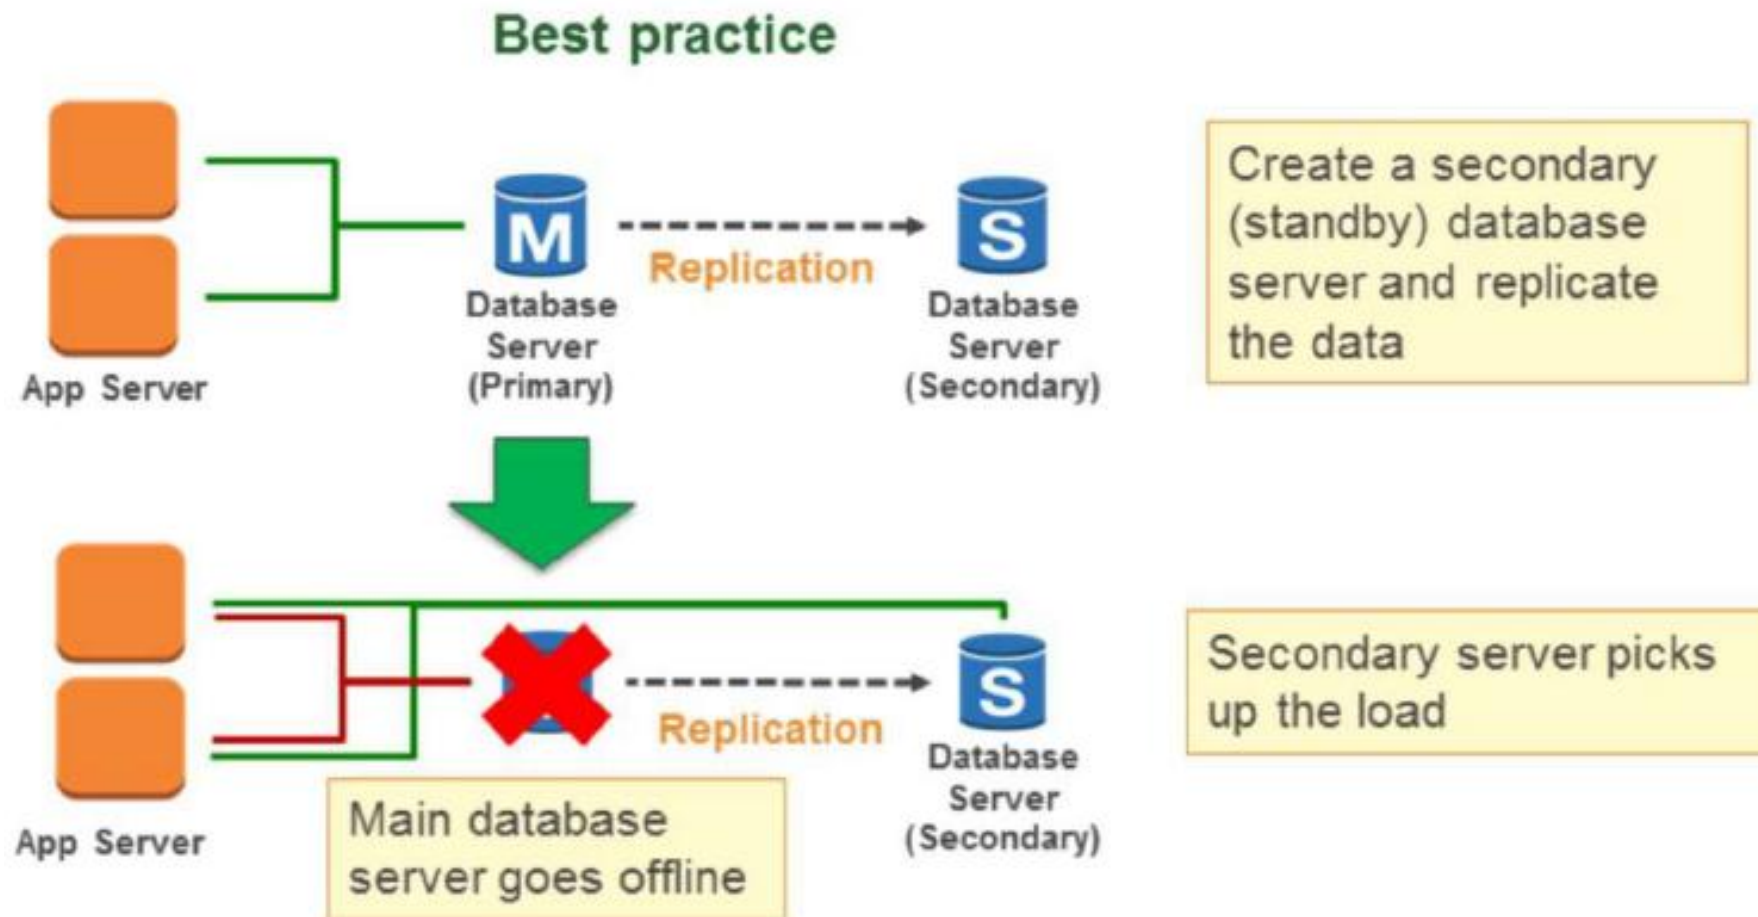
Assume everything fails, and design backwards.

Implement redundancy where possible in order to prevent single failures from bringing down an entire system.



- It does not mean you have to duplicate the resources
- More important is recovery automation and the use of services that provide HA

Best Practice: Avoid Single Points Of Failure



High Availability: Driven By Your Requirements


📦 Recovery Time Objective (**RTO**)

➤ How quickly must the system recover?

📦 Recovery Point Objective (**RPO**)

➤ How much data can you afford to lose?

How much money do you need to invest to meet those objectives?



High Availability Factors

❏ **Fault tolerance:**

The **built-in redundancy** of an application's components.

❏ **Recoverability:**

The process, policies, and procedures related to **restoring service** after a catastrophic event.

❏ **Scalability:**

The ability of an application to **accommodate growth** without changing design.

High Availability and AWS Regions

Multi-region deployment increases:

- Availability
- Cost
- Complexity



Default to a **single region**, unless multi-region deployment is necessary.
If you use a single region, **multi-AZ is the bare minimum** HA solution.

AWS Services and High Availability

Inherently HA services

- ✓ Amazon S3 and Amazon Glacier
- ✓ DynamoDB
- ✓ Amazon CloudFront
- ✓ Amazon SWF
- ✓ Amazon SQS
- ✓ Amazon SNS
- ✓ Amazon SES
- ✓ Amazon Route 53
- ✓ Elastic Load Balancing
- ✓ IAM
- ✓ Amazon CloudWatch
- ✓ AWS Data Pipeline
- ✓ Amazon Kinesis
- ✓ Auto Scaling
- ✓ Amazon Elastic File System
- ✓ AWS CloudFormation
- ✓ Amazon WorkMail
- ✓ AWS Directory Service
- ✓ AWS Lambda
- ✓ Amazon EBS
- ✓ Amazon RDS

HA with the right architecture

- ✦ Amazon EC2
- ✦ Amazon VPC
- ✦ Amazon Redshift
- ✦ Amazon ElastiCache
- ✦ AWS Direct Connect

****Not all services are listed here.**

ELB – Elastic Load Balancer



A managed load balancing service that **distributes incoming application traffic** across multiple Amazon EC2 instances.

AWS Load Balancing Solutions - ELB

Elastic Load Balancing:

- ❏ Distributes load between instances.
- ❏ Recognizes and responds to unhealthy instances.
- ❏ Can be public or internal-facing.
- ❏ Uses HTTP, HTTPS, and TCP protocols.
- ❏ Each load balancer is given a public DNS name.
 - **Internet-facing** load balancers have DNS names which publicly resolve to the public IP addresses of the load balancer's nodes.
 - **Internal** load balancers have DNS names which publicly resolve to the private IP addresses of the load balancer's nodes.

Reasons for Using ELB

Elastic Load Balancing provides the following features:

- 📦 Health checks
- 📦 Cross-zone load balancing
- 📦 Proxy Protocol
- 📦 Sticky sessions
- 📦 Connection draining



ELB – Connection Draining

Enabling *connection draining* causes the load balancer to **stop sending new requests** to the back-end instances when instances are **de-registering** or **become unhealthy**.

Why is this important?

You can perform maintenance without affecting your end users.

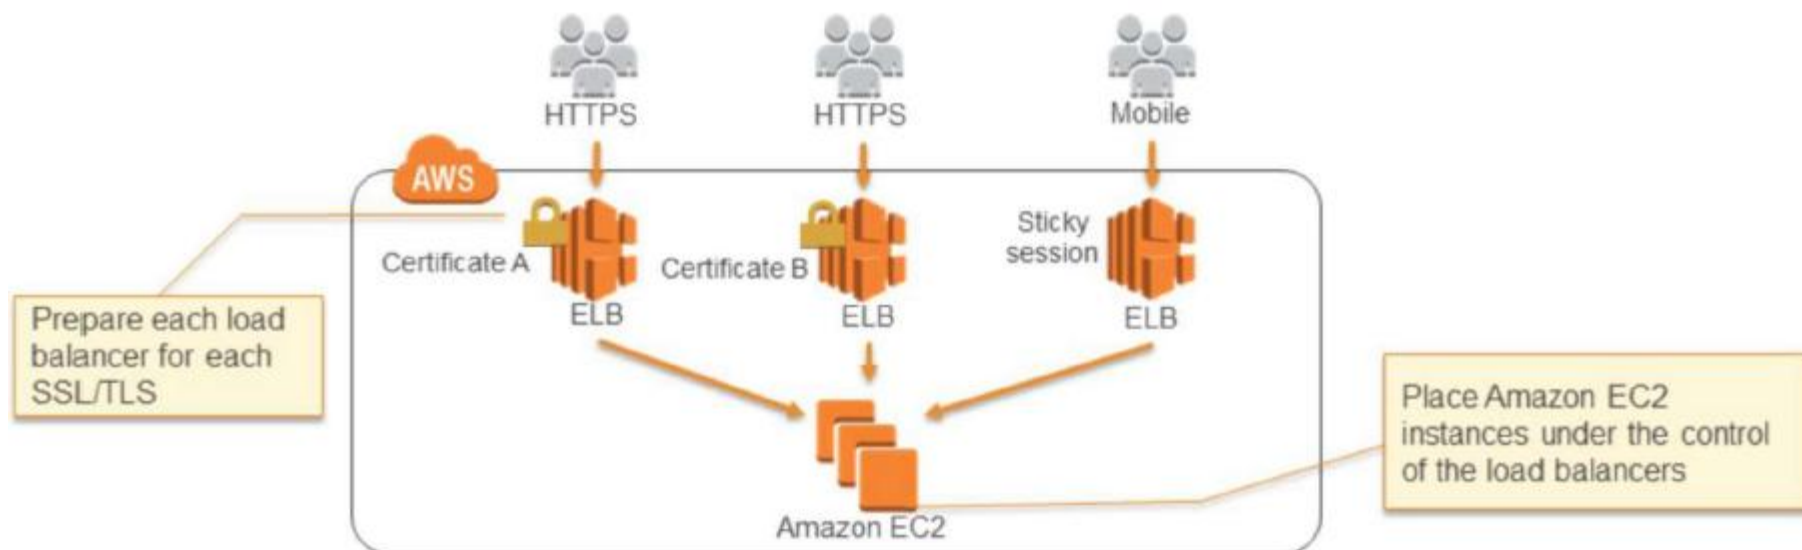
Cloud Design Pattern: Multi Load Balancer Pattern

Problem to be solved:

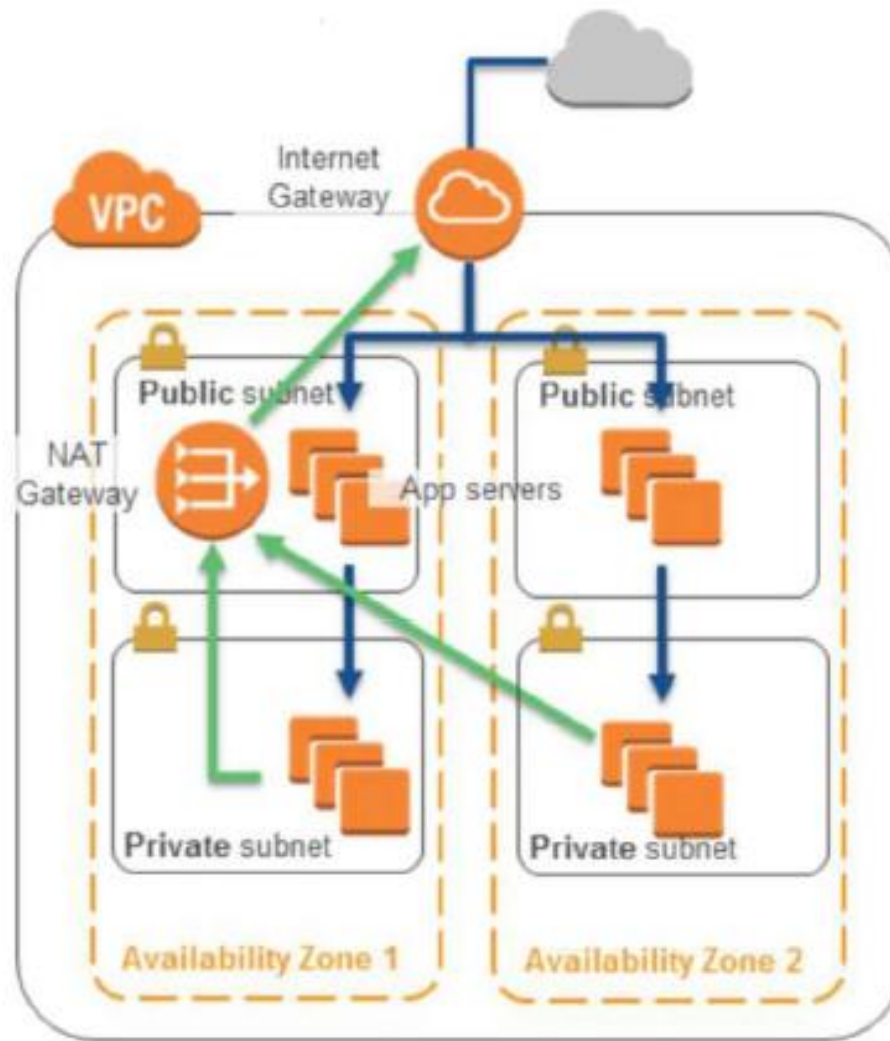
When a web application is multi-device compatible, there will be access from PCs, mobile phones, and smart phones. At this time, if it is necessary to perform a setup such as for SSL/TLS or to assign sessions for individual access devices, and if the setup is performed by the EC2 instances themselves, any change to the settings would become extremely laborious as the number of servers increases.

Cloud Pattern:

Assign multiple load balancers with different settings for the types of devices that are accessing the web application.

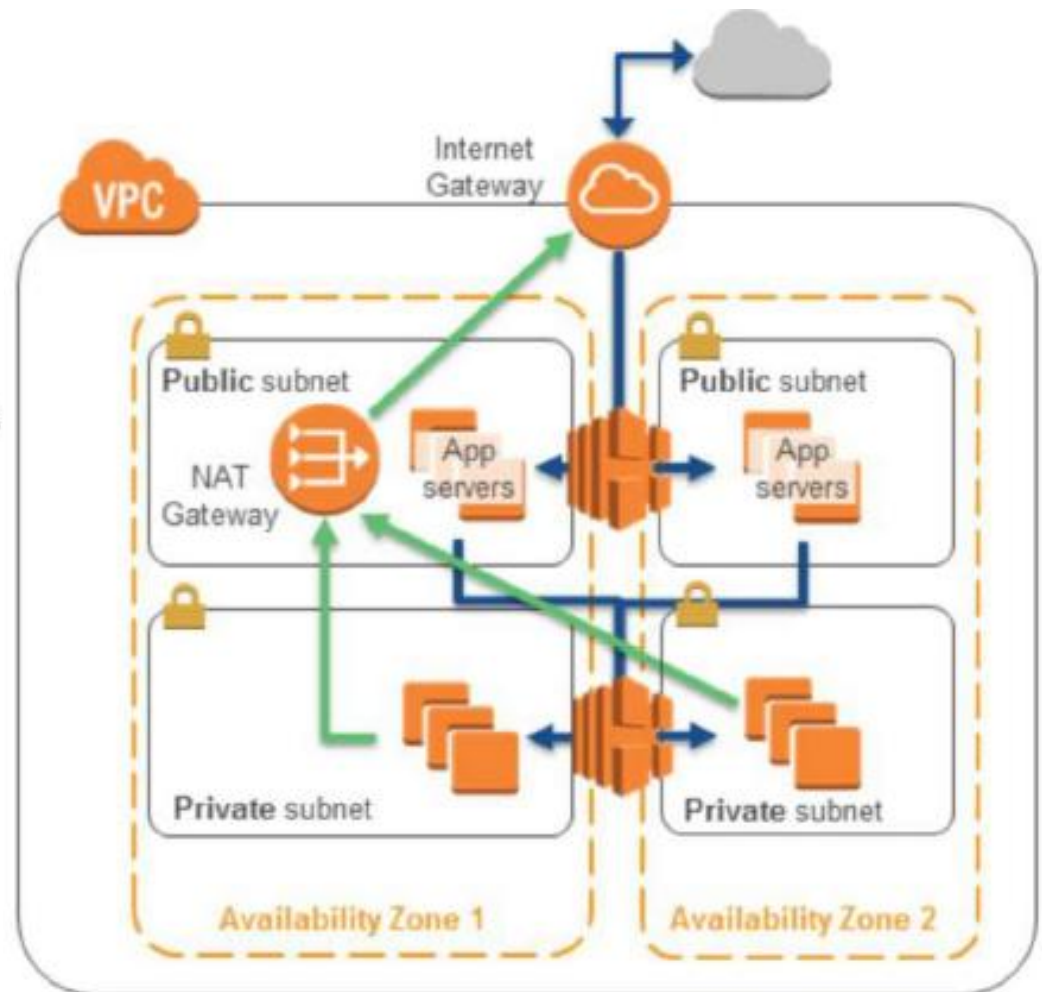


Exercise: How Can Availability of This Environment Be Improved?



High Availability With Load Balancing

- 📦 **Internet-facing** load balancer
 - Deployed in **public** subnet
 - Balancing traffic to app servers in two Availability Zones
- 📦 **Internal** load balancer
 - Intranet-facing or internal




Elastic IP: For Greater Fault Tolerance



Elastic IP Addresses:

- 📦 Are **static IP addresses** designed for dynamic cloud computing.
- 📦 Can be attached to **Amazon EC2 instances**.
- 📦 Enable you to **mask the failure** of an instance or software by allowing your users and clients to **use the same IP address with replacement resources**.

Using Elastic IP Addresses to Enable High Availability



If one instance goes down, clients can use the same IP address to reach the replacement instance.

High Availability Across Regions via Route 53

Route 53 –Highly Available Domain Name Service (DNS)

Amazon Route 53 is an authoritative DNS service from AWS with the following characteristics:

- 📁 DNS translates domain names (like www.amazon.com) into IP addresses .

The name refers to the fact that DNS servers respond to queries on port 53.

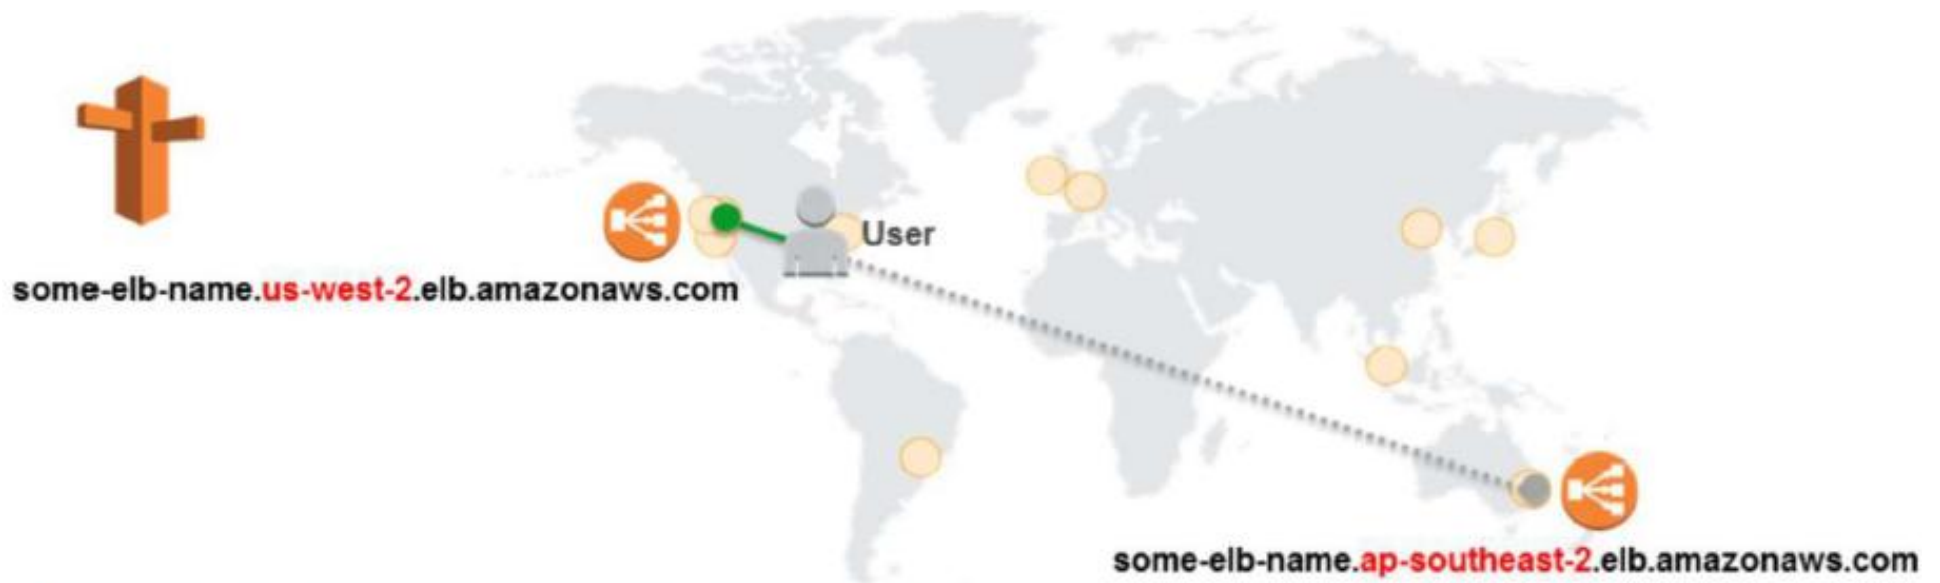


Amazon
Route 53

AWS Route 53 – Routing Options

- ❏ **Simple routing:** Single server environments.
- ❏ **Weighted round robin:** Assign weights to resource record sets to specify the frequency.
- ❏ **Latency-based routing:** Helps to improve your global applications.
- ❏ **Health check and DNS failover:** Fail over to a backup site if your primary site becomes unreachable.
- ❏ **Geolocation routing:** Specify **geographic locations** by continent, by country, or by state in the United States.

Use Case: Multi-Region Deployment



Name	Type	Value
amgogreen.com	ALIAS	some-elb-name.us-west-2.elb.amazonaws.com
amgogreen.com	ALIAS	some-elb-name.ap-southeast-2.elb.amazonaws.com

AWS High Availability

Improve availability of your applications running on AWS by:

- 📦 Configuring backup and failover scenarios for your own applications.
- 📦 Enabling highly available multi-region architectures on AWS.
- 📦 Improving health checks by combining multiple health checks, domain name–based health checks, string matching, specifying the request interval, and more.

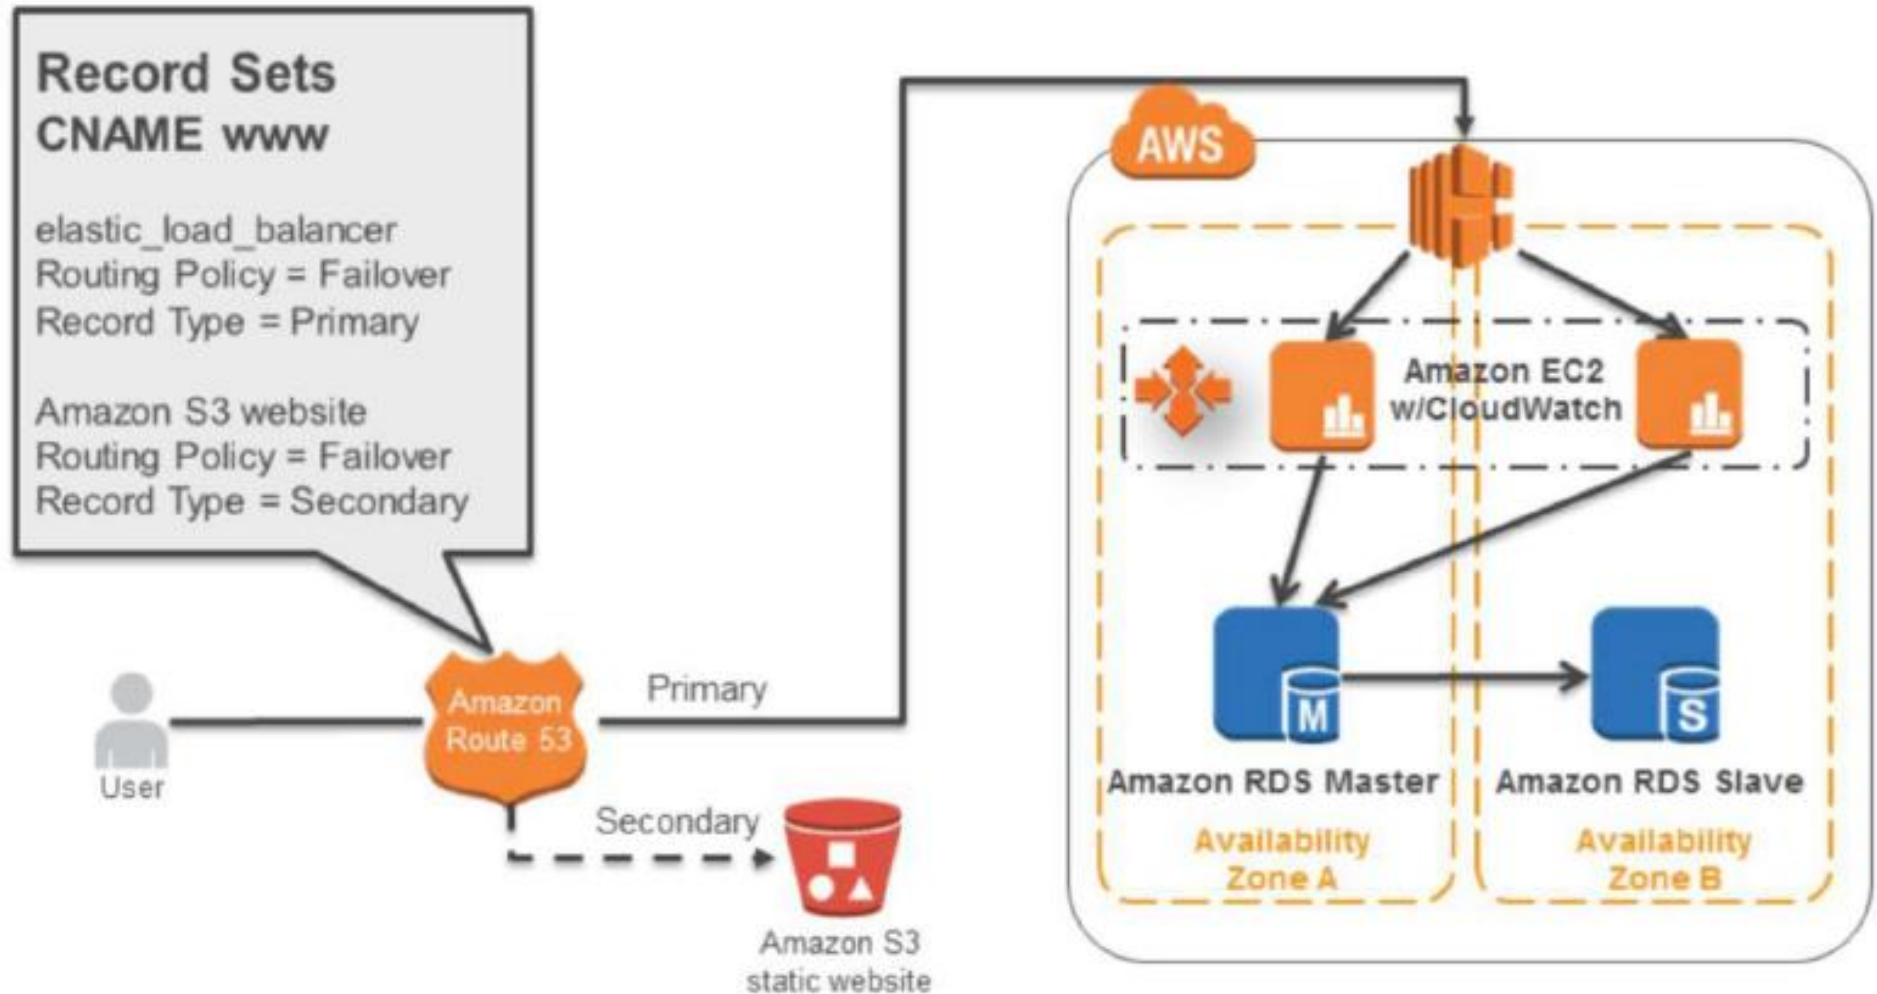
📦 *Solution is driven by your SLA*

📦 *Robust HA solutions are mainly determined by the **perfection** of the **recovery chain**:*

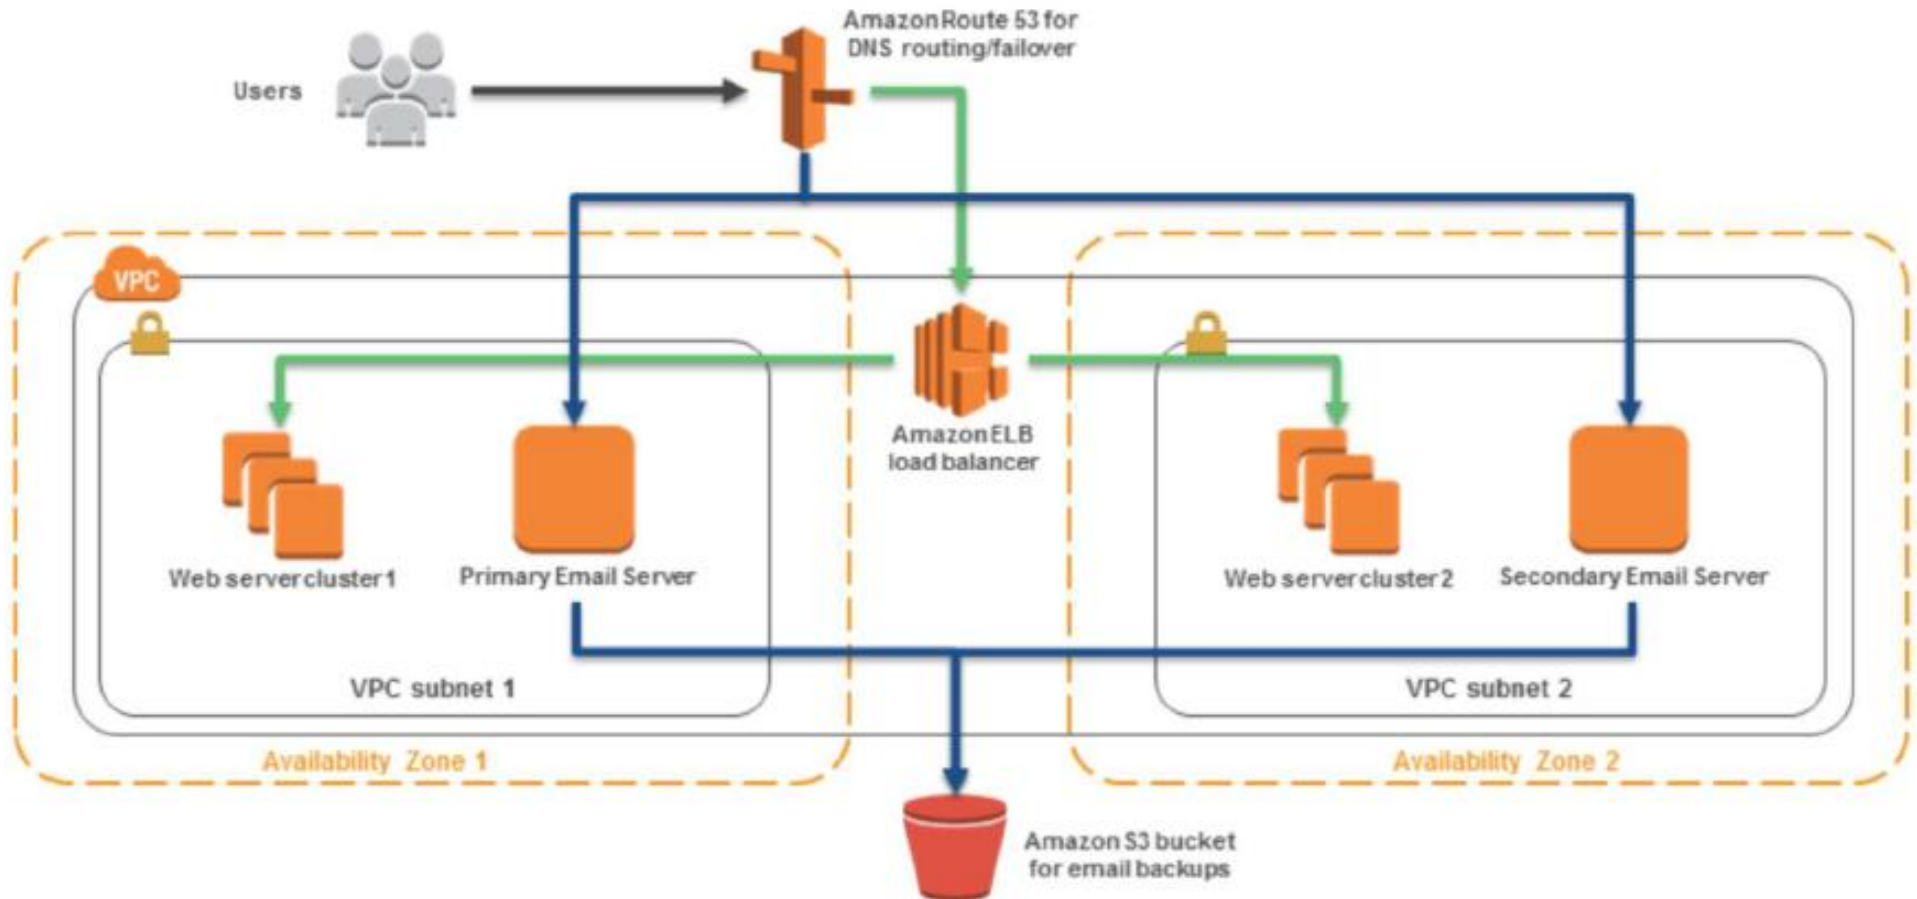
- *failure detection*
- *flawless failover code*
- *working backups*

<http://ucare.cs.uchicago.edu/pdf/socc16-cos.pdf>

Typical Multi Region Architecture



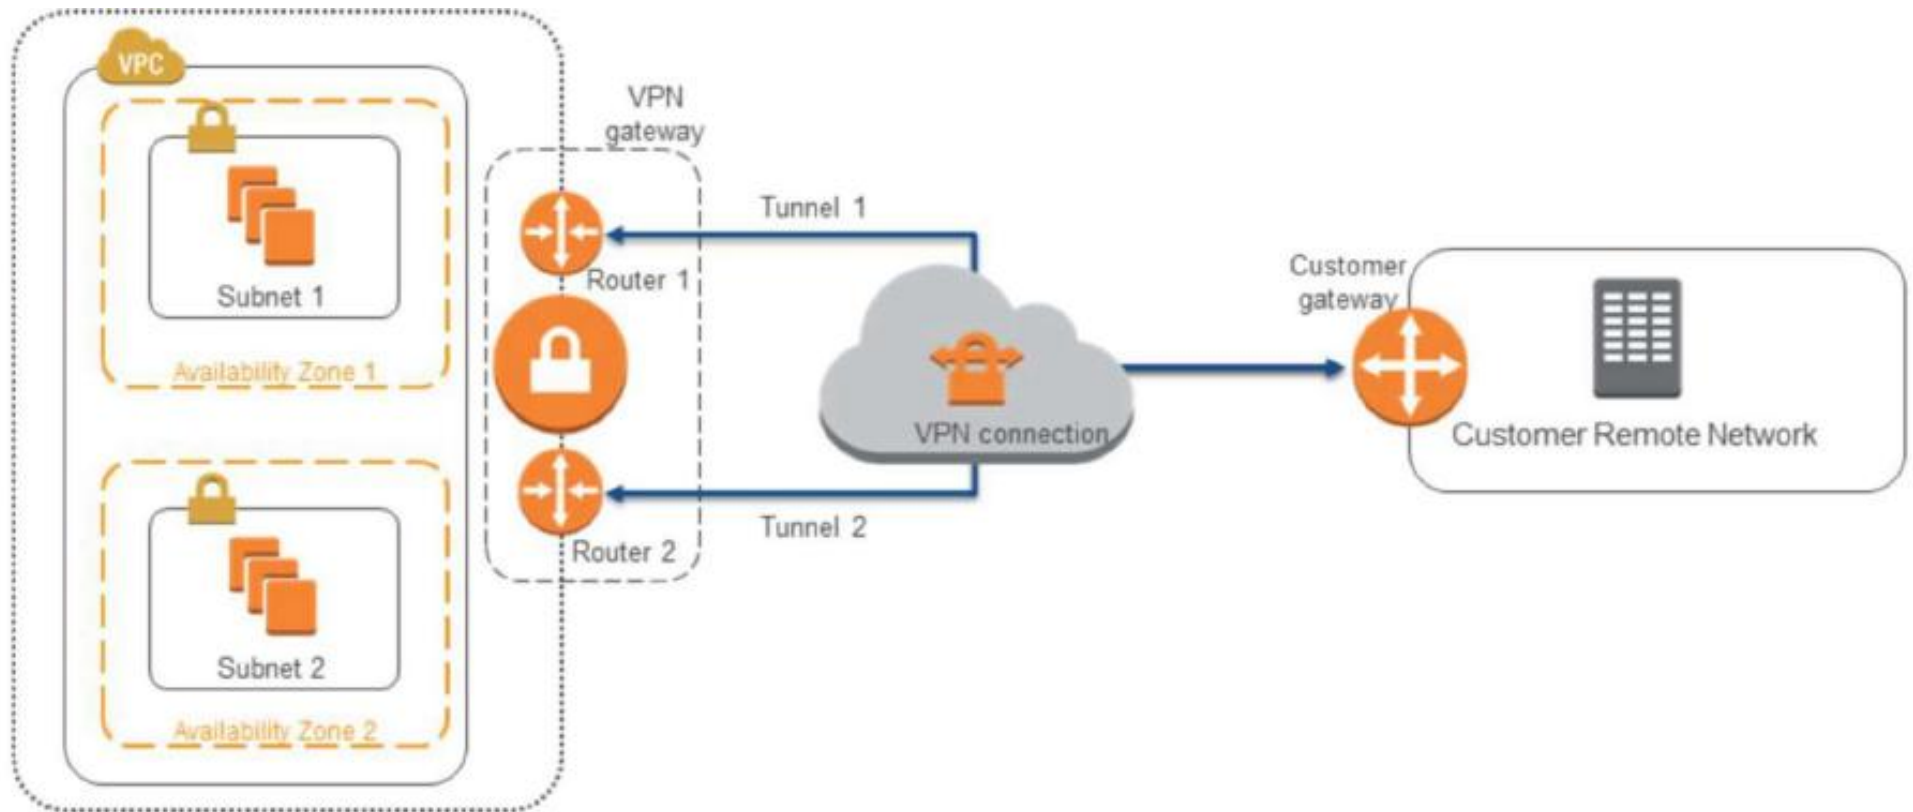
Alta Veiculos Use Case



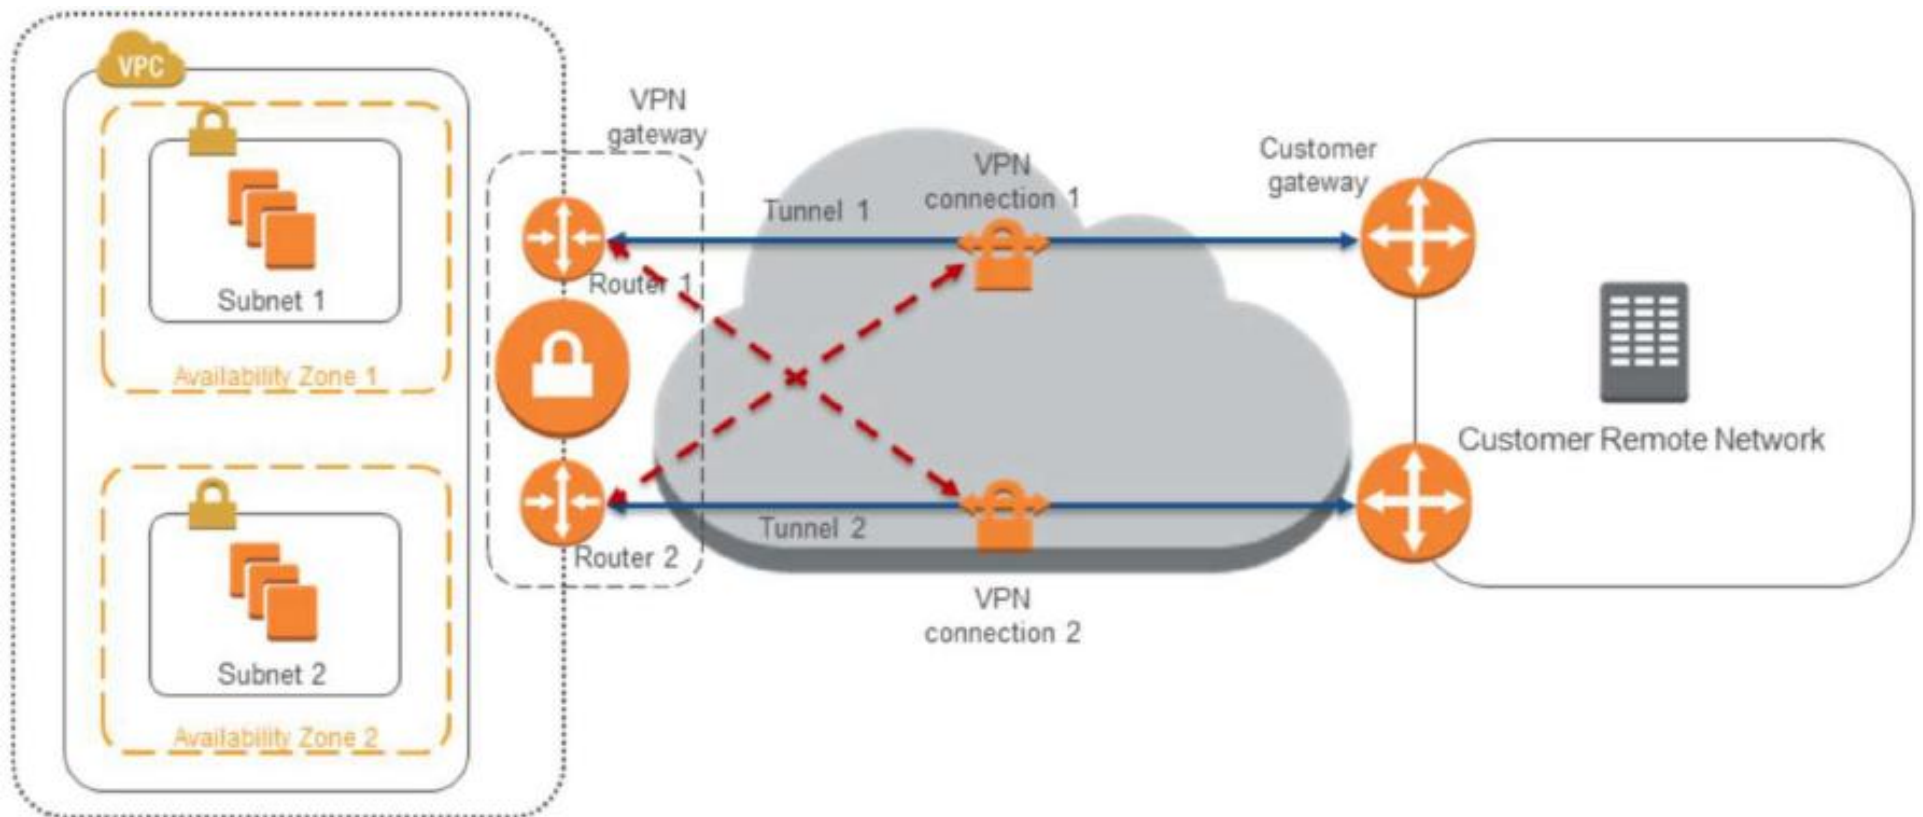
For more on this case study, see: <http://aws.typepad.com/brasil/case-studies/>

Availability Considerations About Connections to Components Outside of AWS

VPN Connection : Is This Highly Available?



Adding Connection Redundancy With AWS HW VPNs



Best Practice : Enable Scalability

Ensure that your architecture can handle changes in demand.


A key advantage of a cloud-based infrastructure is how **quickly** you can respond to changes in resource needs.



Best Practice : Enable Scalability

Ensure that your architecture can handle changes in demand.

A key advantage of a cloud-based infrastructure is how **quickly** you can respond to changes in resource needs.



Vertical vs Horizontal Scaling

Vertical scaling

Scale up and down


- Change in the **specifications** of instances (more CPU, memory, etc.)



Horizontal scaling

Scale in and out

- Change in the **number** of instances (Add and remove instances as needed)



What Information Is Needed For Scaling

To leverage AWS in an **efficient** way, you need **insight** into your AWS resources, e.g.:

- 📦 How much of your infrastructure is actually being used?
- 📦 Is your application's performance or availability being affected by a lack of sufficient capacity?

How do we collect this information ??

Collecting Metrics: AWS Cloud Watch



Amazon CloudWatch:

- 📦 **Monitors components** in your infrastructure based on what you identify.
- 📦 **Sends notifications and triggers auto scaling** actions based on metrics you specify.

Monitor AWS Resources With Cloud Watch

Amazon CloudWatch offers:

- 📦 A distributed statistics gathering system; it collects and tracks metrics.
- 📦 Metrics that are seamlessly collected at the hypervisor level by default.
- 📦 **The ability to create and use custom metrics** of data generated by your own applications and services.

Examples:

- The time web pages take to load
- Request error rates
- Number of simultaneous processes or threads

Cloud Watch Alarm Examples

Amazon
EC2



If CPU utilization is > 60% for 5 minutes...

Amazon
RDS



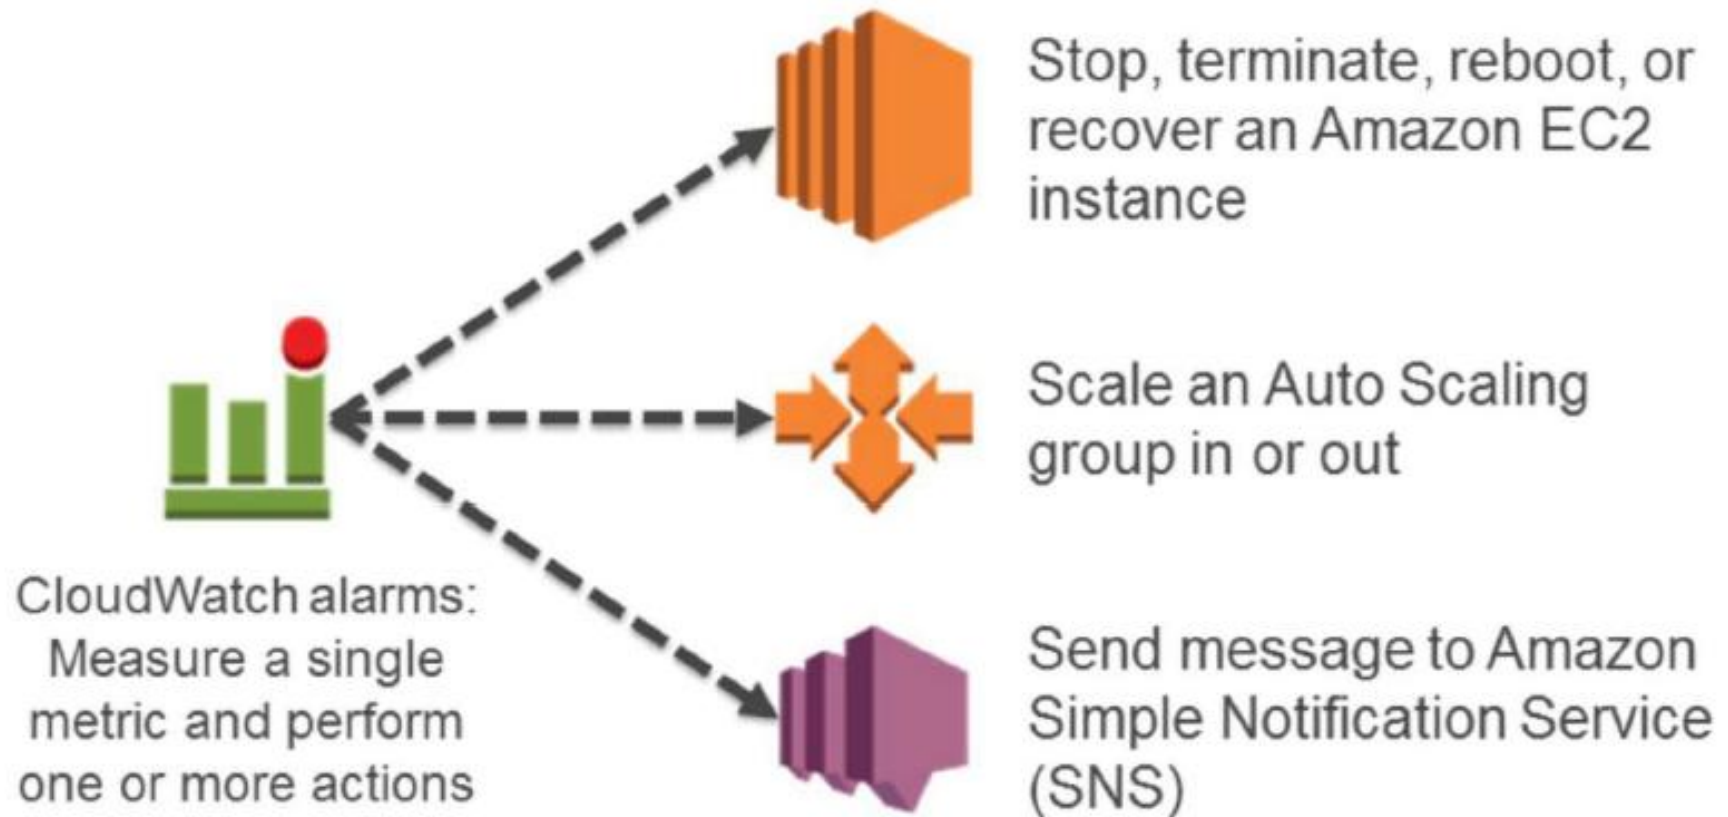
If number of simultaneous connections is > 10 for one minute...

Amazon
ELB



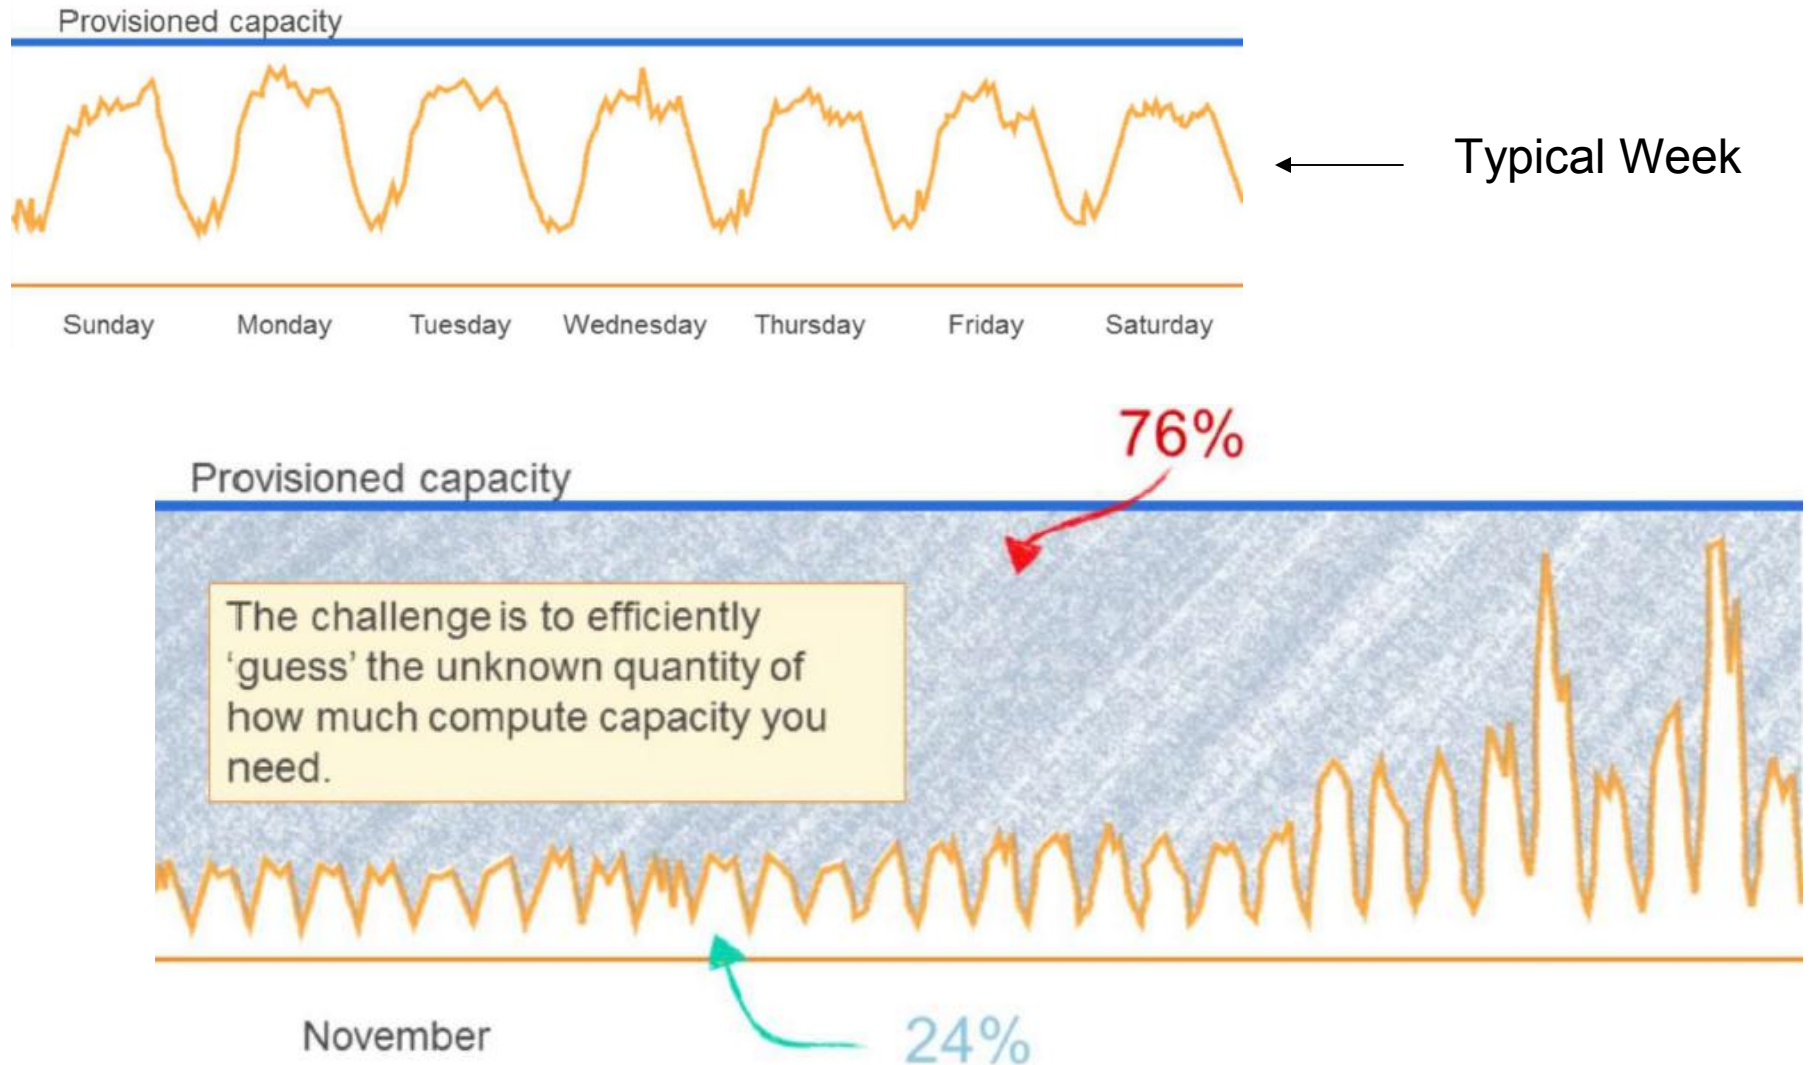
If number of healthy hosts is < 5 for 10 minutes...

Cloud Watch Alarms and Actions




Autoscaling

Traffic Profile For A Major E-Commerce Site



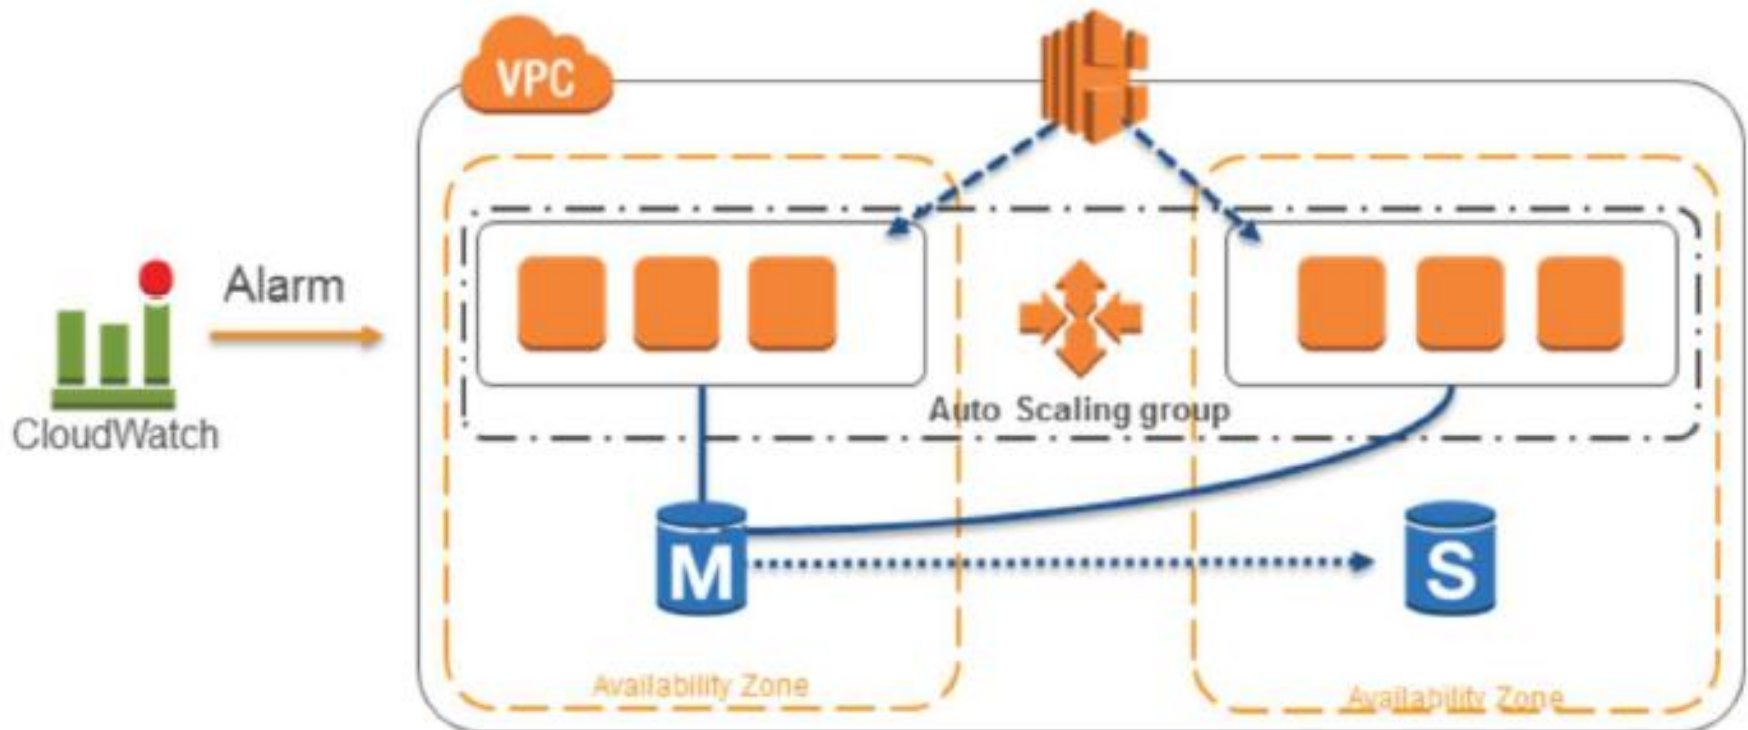
Auto Scaling



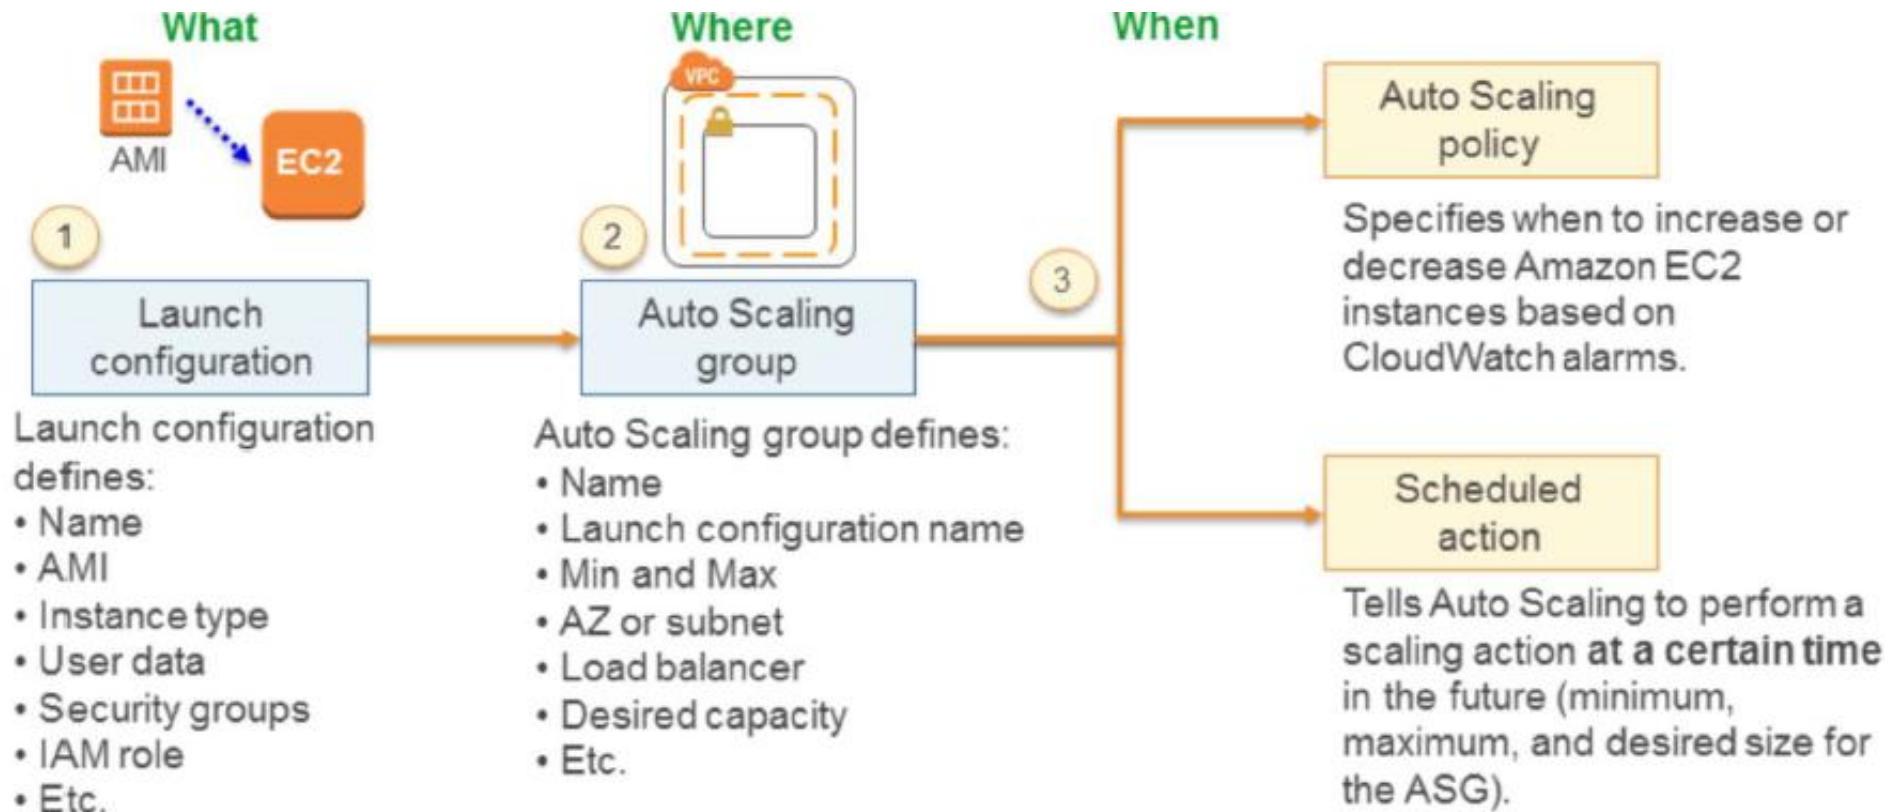
- 📦 Launches or terminates instances based on specified conditions.
- 📦 Automatically registers new instances with load balancers when specified.
- 📦 Can launch across Availability Zones.

No More Guessing About Resource Needs


Build a **flexible** system that will react to changes in customer demand and manage costs dynamically.



How Does Auto Scaling Work?



ELB + AS + CW: More Power Putting Connecting The Services



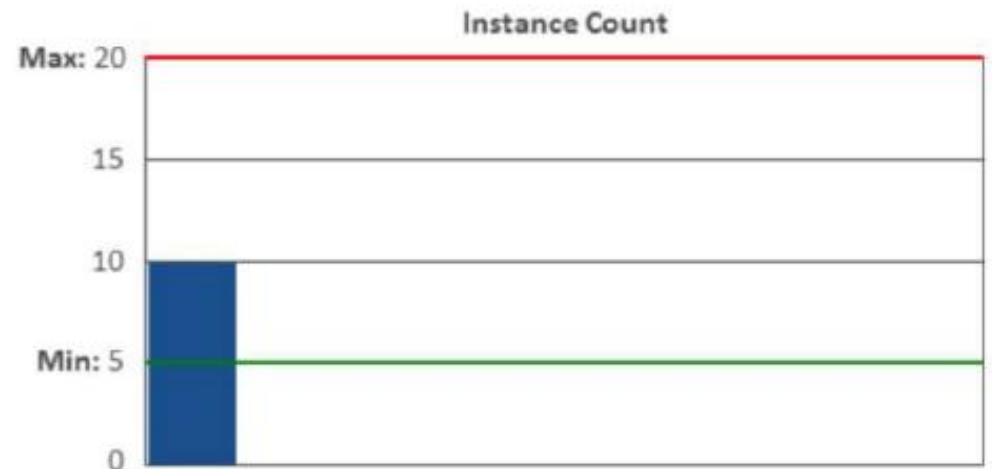

Auto Scaling Steps

Alarms:

- 🔔 Add 2 instances when average CPU is 80-100%
- 🔔 Add 1 instance when average CPU is 60-80%
- 🔔 Remove 1 instance when average CPU is 20-40%
- 🔔 Remove 2 instances when average CPU is 0-20%

Limits:

- 🔔 Minimum: 5 instances
- 🔔 Maximum: 20 instances




Auto Scaling Steps

As usage increases:

- 📦 CPU utilization goes up.

When CPU utilization is 60-80%:

- 📦 Scale-out alarm is triggered.
- 📦 Add 1 step policy is applied.
- 📦 New instance is launched but not added to the aggregated group metrics until after warm up period expires.




Auto Scaling Steps

As usage increases:

- 📊 CPU utilization goes up.

While waiting for new instance:

- 📊 CPU utilization remains high.
- 📊 Another alarm period is triggered.
- 📊 Since current capacity is still 10 during the warmup period, and desired capacity is already 11, **no additional instances are launched.**




Auto Scaling Steps

As usage increases further:

- 📦 CPU utilization goes up.

When CPU utilization is 80-100%:

- 📦 Scale-out alarm is triggered.
- 📦 Add 2 step policy is applied.
- 📦 Since the alarm occurred during a warm up period, two instances are launched less the one instance added during the first alarm.
- 📦 Again new instances are not added to aggregated group metrics.




Auto Scaling Steps

As capacity matches usage:

- 📦 CPU utilization stabilizes.

When CPU utilization is 40-60%:

- 📦 No alarms are triggered.
- 📦 After warm up periods expire, new instances are added to the aggregated group metrics.




Auto Scaling Steps

As usage decreases:

- 📦 CPU utilization goes down.

When CPU utilization is 20-40%:

- 📦 Scale-in alarm is triggered.
- 📦 Remove 1 step policy is applied.
- 📦 An instance is removed from the Auto Scaling group and from the aggregated group metrics.




Auto Scaling Steps

As usage decreases:

- 📦 CPU utilization goes down further.

When CPU utilization is 0-20%:

- 📦 Scale in alarm is triggered.
- 📦 Remove 2 step policy is applied.
- 📦 Two instances are removed from the Auto Scaling group and from the aggregated group metrics.





Auto Scaling Steps

As capacity matches usage:

- 📦 CPU utilization stabilizes.

When $40\% < \text{CPU Utilization} < 60\%$

- 📦 No alarm is triggered.
- 📦 No step policies are applied.
- 📦 No instances are added or removed from service.

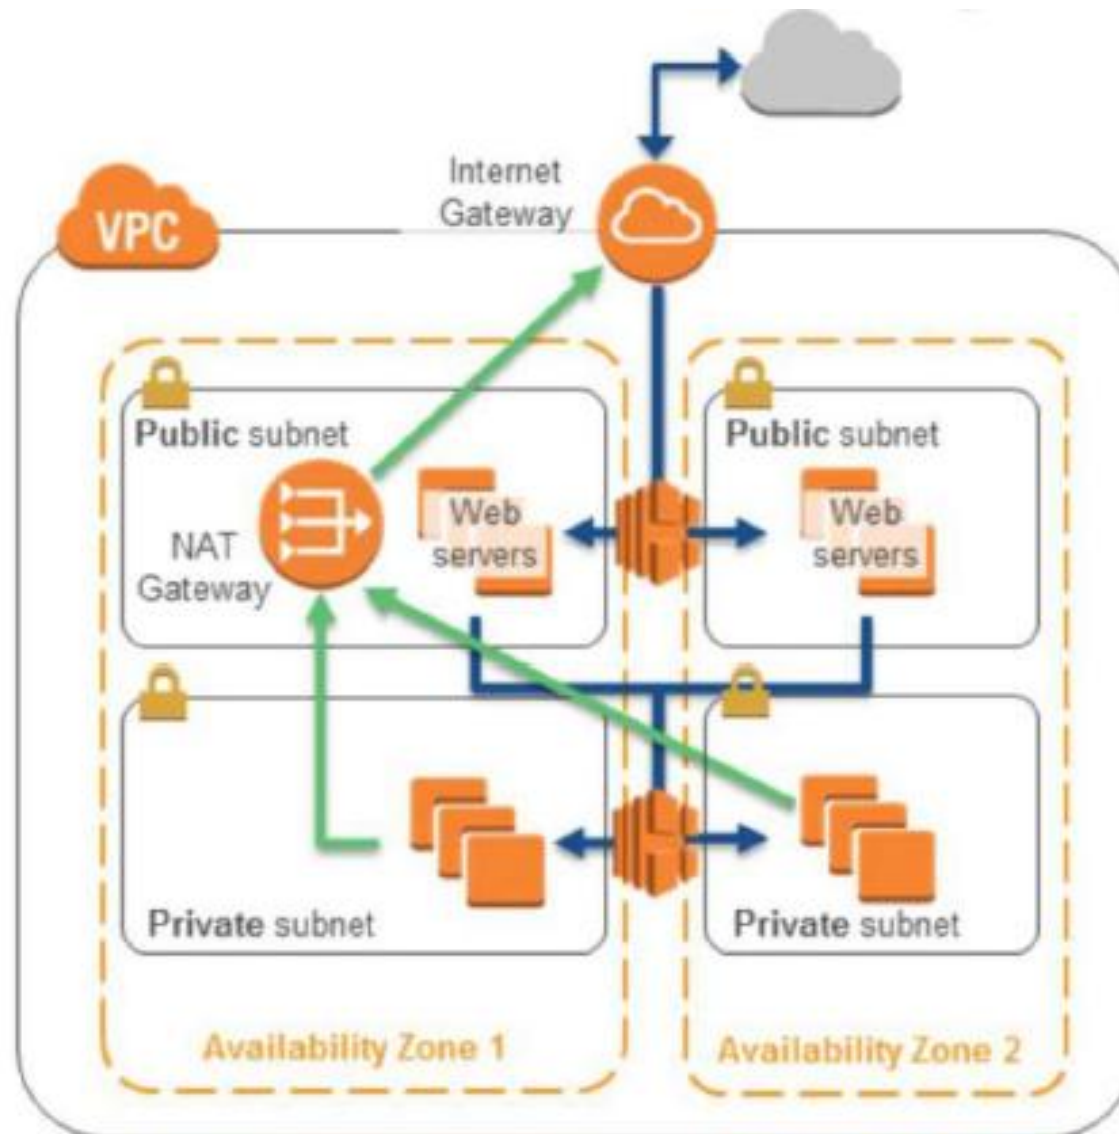


Auto Scaling Considerations

- ❏ Avoid Auto Scaling thrashing.
 - Be more cautious about scaling in; avoid aggressive instance termination.
 - Scale out early, scale in slowly.
- ❏ Set the min and max capacity parameter values carefully.
- ❏ Use lifecycle hooks.
 - Perform custom actions as Auto Scaling launches or terminates instances.
- ❏ Stateful applications will require additional automatic configuration of instances launched into Auto Scaling groups.

Remember: Instances can take several minutes after launch to be **fully usable**.

How Can You Increase Availability Of This Configuration?



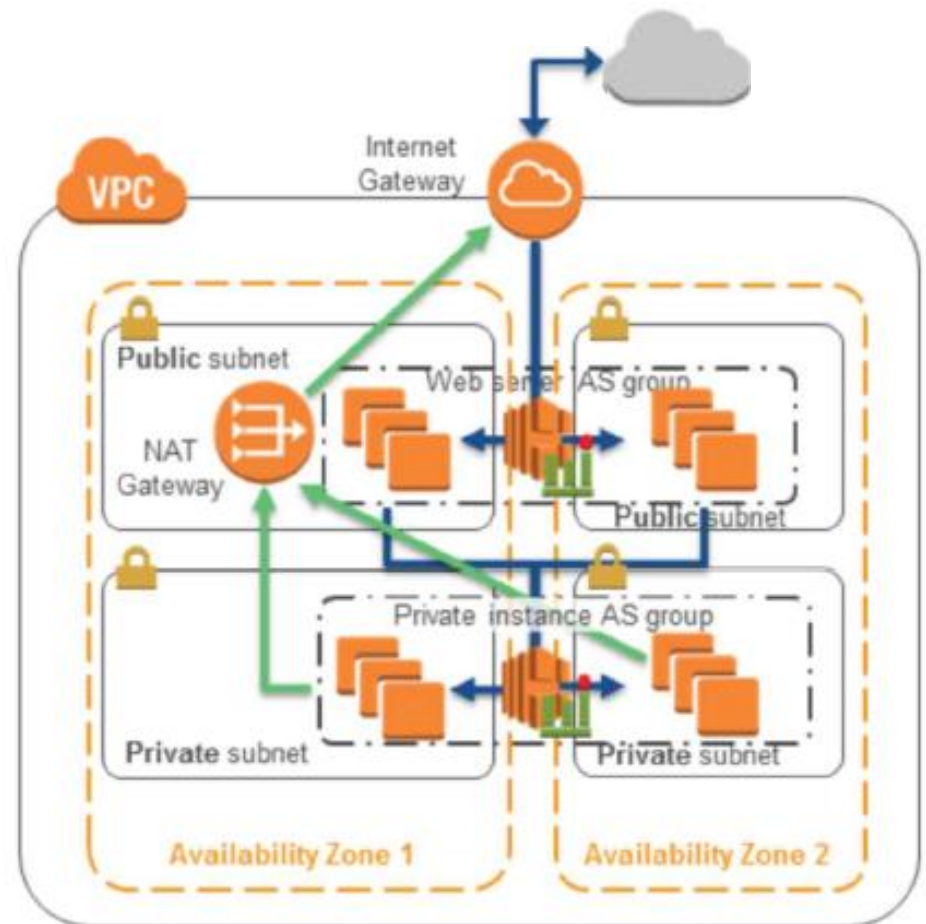
HA With Auto Scaling and Cloud Watch

Amazon CloudWatch

- 📦 **Monitors** the health of your ELBs and/or your instances and **launches actions** appropriately.

Auto Scaling

- 📦 Receives request from **Amazon CloudWatch**.
- 📦 **Adds or subtracts** instances.



AWS Lambda and Event Driven Scaling



- ❏ Fully managed compute service that runs stateless code (Node.js, Java, and Python) in response to an event or on a time-based interval.
- ❏ Allows you to run code without managing infrastructure like Amazon EC2 instances and Auto Scaling groups.

AWS Lambda Service

AWS Lambda **handles**:

- 📦 Servers
- 📦 Capacity needs
- 📦 Deployment
- 📦 Scaling and fault tolerance
- 📦 OS or language updates
- 📦 Metrics and logging

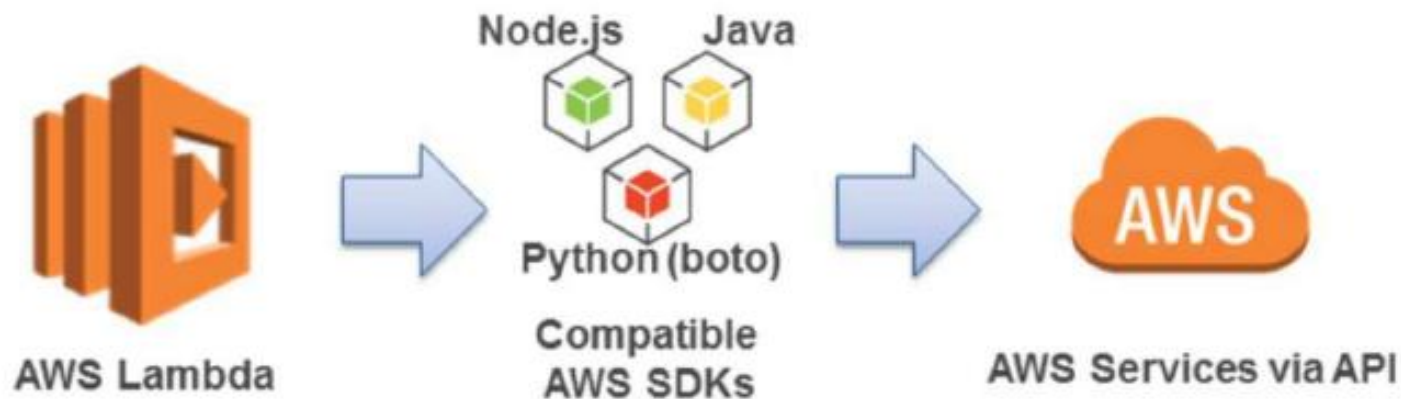
AWS Lambda **enables** you to:

- 📦 Bring your own code (even native libraries).
- 📦 Run code in parallel.
- 📦 Create back ends, event handlers, and data processing systems.
- 📦 Never pay for idling resources!

How AWS Lambda Can Be Used With Scaling?

Scaling events can trigger AWS Lambda functions. In addition, code run in AWS Lambda has access to the AWS API and can grant permissions via AWS IAM roles.

This means you can use a Lambda function to **automatically make API calls to other AWS services** when scaling happens.



How AWS Lambda Can Be Used With Scaling?

Examples of scaling-related operations you could perform with AWS Lambda:

- 📦 Scale container-based instances (Docker, Amazon Elastic Container Service, etc.)
- 📦 Scale more intelligently using functions (e.g.: analyze a stream of performance data looking for patterns rather than just events)
- 📦 Since AWS Lambda can scale automatically, consider replacing some Amazon EC2 instances with Lambda functions where appropriate

AWS Lambda and Event Driven Scaling



- ❏ Fully managed compute service that runs stateless code (Node.js, Java, and Python) in response to an event or on a time-based interval.
- ❏ Allows you to run code without managing infrastructure like Amazon EC2 instances and Auto Scaling groups.

AWS Lambda and Event Driven Scaling



- ❏ Fully managed compute service that runs stateless code (Node.js, Java, and Python) in response to an event or on a time-based interval.
- ❏ Allows you to run code without managing infrastructure like Amazon EC2 instances and Auto Scaling groups.

Resources

- AWS Best Practices:
https://d0.awsstatic.com/whitepapers/AWS_Cloud_Best_Practices.pdf
- AWS Well Architected Framework:
https://d0.awsstatic.com/whitepapers/architecture/AWS_Well-Architected_Framework.pdf
- AWS SlideShare Channel:
<https://www.slideshare.net/AmazonWebServices/>